



In [ ]:

```
!pip install transformers==2.4.0
```

Collecting transformers==2.4.0

Downloading <https://files.pythonhosted.org/packages/c6/38/c30b6a4b86705311c428a234ef752f6c4c4ffdd75422a829f1f2766136c3/transformers-2.4.0-py3-none-any.whl> (475kB)

|██| 481kB 3.4MB/s eta 0:00:01

Requirement already satisfied: boto3 in /usr/local/lib/python3.6/dist-packages (from transformers==2.4.0) (1.14.22)

Requirement already satisfied: filelock in /usr/local/lib/python3.6/dist-packages (from transformers==2.4.0) (3.0.12)

Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.6/dist-packages (from transformers==2.4.0) (2019.12.20)

Collecting sentencepiece

Downloading [https://files.pythonhosted.org/packages/d4/a4/d0a884c4300004a78cca907a6ff9a5e9fe4f090f5d95ab341c53d28cbc58/sentencepiece-0.1.91-cp36-cp36m-manylinux1\\_x86\\_64.whl](https://files.pythonhosted.org/packages/d4/a4/d0a884c4300004a78cca907a6ff9a5e9fe4f090f5d95ab341c53d28cbc58/sentencepiece-0.1.91-cp36-cp36m-manylinux1_x86_64.whl) (1.1MB)

|██| 1.1MB 34.7MB/s

Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.6/dist-packages (from transformers==2.4.0) (4.41.1)

Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from transformers==2.4.0) (1.18.5)

Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages (from transformers==2.4.0) (2.23.0)

Collecting tokenizers==0.0.11

Downloading [https://files.pythonhosted.org/packages/5e/36/7af38d572c935f8e0462ec7b4f7a46d73a2b3b1a938f50a5e8132d5b2dc5/tokenizers-0.0.11-cp36-cp36m-manylinux1\\_x86\\_64.whl](https://files.pythonhosted.org/packages/5e/36/7af38d572c935f8e0462ec7b4f7a46d73a2b3b1a938f50a5e8132d5b2dc5/tokenizers-0.0.11-cp36-cp36m-manylinux1_x86_64.whl) (3.1MB)

|██| 3.1MB 31.5MB/s

Collecting sacremoses

Downloading <https://files.pythonhosted.org/packages/7d/34/09d19aff26edcc8eb2a01bed8e98f13a1537005d31e95233fd48216eed10/sacremoses-0.0.43.tar.gz> (883kB)

|██| 890kB 34.3MB/s

Requirement already satisfied: s3transfer<0.4.0,>=0.3.0 in /usr/local/lib/python3.6/dist-packages (from boto3->transformers==2.4.0) (0.3.3)

Requirement already satisfied: jmespath<1.0.0,>=0.7.1 in /usr/local/lib/python3.6/dist-packages (from boto3->transformers==2.4.0) (0.10.0)

Requirement already satisfied: botocore<1.18.0,>=1.17.22 in /usr/local/lib/python3.6/dist-packages (from boto3->transformers==2.4.0) (1.17.22)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.6/dist-packages (from requests->transformers==2.4.0) (2020.6.20)

Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.6/dist-packages (from requests->transformers==2.4.0) (1.24.3)

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.6/dist-packages (from requests->transformers==2.4.0) (2.10)

Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.6/dist-packages (from requests->transformers==2.4.0) (3.0.4)

Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from sacremoses->trans

```

formers==2.4.0) (1.15.0)
Requirement already satisfied: click in /usr/local/lib/python3.6/dist-packages (from sacremoses->transformers==2.4.0) (7.1.2)
Requirement already satisfied: joblib in /usr/local/lib/python3.6/dist-packages (from sacremoses->transformers==2.4.0) (0.16.0)
Requirement already satisfied: docutils<0.16,>=0.10 in /usr/local/lib/python3.6/dist-packages (from botocore<1.18.0,>=1.17.22->boto3->transformers==2.4.0) (0.15.2)
Requirement already satisfied: python-dateutil<3.0.0,>=2.1 in /usr/local/lib/python3.6/dist-packages (from botocore<1.18.0,>=1.17.22->boto3->transformers==2.4.0) (2.8.1)
Building wheels for collected packages: sacremoses
  Building wheel for sacremoses (setup.py) ... done
  Created wheel for sacremoses: filename=sacremoses-0.0.43-cp36-none-any.whl size=893260 sha256=d568057d58e8e6b796eaab43d560028cff8576d82b291abbd575f7eedfe61870
  Stored in directory: /root/.cache/pip/wheels/29/3c/fd/7ce5c3f0666dab31a50123635e6fb5e19ceb42ce38d4e58f45
Successfully built sacremoses
Installing collected packages: sentencepiece, tokenizers, sacremoses, transformers
Successfully installed sacremoses-0.0.43 sentencepiece-0.1.91 tokenizers-0.0.11 transformers-2.4.0

```

In [ ]:

```

# importing necessary libraries
from typing import List, Tuple
import random
import html

import pandas as pd
import numpy as np
from sklearn.model_selection import GroupKFold, KFold
import matplotlib.pyplot as plt
from tqdm.notebook import tqdm
import tensorflow as tf
import tensorflow.keras.backend as K
import os
from scipy.stats import spearmanr
from scipy.optimize import minimize
from math import floor, ceil
from transformers import *
from tensorflow.keras.layers import Flatten, Dense, Dropout, GlobalAveragePooling1D
from tensorflow.keras.models import Model

```

In [ ]:

```
# fixing random seeds
seed = 13
random.seed(seed)
os.environ[ 'PYTHONHASHSEED' ] = str(seed)
np.random.seed(seed)
tf.random.set_seed(seed)
```

In [ ]:

```
# from google.colab import drive
# drive.mount('/content/drive')
```

In [ ]:

```
!wget https://github.com/SarthakV7/Kaggle_google_quest_challenge/blob/master/test.csv
!wget https://raw.githubusercontent.com/SarthakV7/Kaggle_google_quest_challenge/master/train.csv
```

```
--2020-07-23 16:19:30-- https://github.com/SarthakV7/Kaggle_google_quest_challenge/blob/master/test.csv
```

```
Resolving github.com (github.com)... 140.82.114.3
```

```
Connecting to github.com (github.com)|140.82.114.3|:443... connected.
```

```
HTTP request sent, awaiting response... 200 OK
```

```
Length: unspecified [text/html]
```

```
Saving to: 'test.csv'
```

```
test.csv          [  <=>          ]   3.87M  15.4MB/s   in 0.3s
```

```
2020-07-23 16:19:31 (15.4 MB/s) - 'test.csv' saved [4053810]
```

```
--2020-07-23 16:19:33-- https://raw.githubusercontent.com/SarthakV7/Kaggle_google_quest_challenge/master/train.csv
```

```
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 151.101.0.133, 151.101.64.133, 151.101.128.133, ...
```

```
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|151.101.0.133|:443... connected.
```

```
HTTP request sent, awaiting response... 200 OK
```

```
Length: 13755514 (13M) [text/plain]
```

```
Saving to: 'train.csv'
```

```
train.csv         100%[=====>]  13.12M  28.5MB/s   in 0.5s
```

```
2020-07-23 16:19:34 (28.5 MB/s) - 'train.csv' saved [13755514/13755514]
```

In [ ]:

```
# reading the data into dataframe using pandas
train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')
submission = pd.read_csv('sample_submission.csv')
```

In [ ]:

```
# Selecting data for training and testing
y = train[train.columns[11:]] # storing the target values in y
X = train[['question_title', 'question_body', 'answer']]
X_test = test[['question_title', 'question_body', 'answer']]
```

In [ ]:

```
# Cleaning the data
X.question_body = X.question_body.apply(html.unescape)
X.question_title = X.question_title.apply(html.unescape)
X.answer = X.answer.apply(html.unescape)

X_test.question_body = X_test.question_body.apply(html.unescape)
X_test.question_title = X_test.question_title.apply(html.unescape)
X_test.answer = X_test.answer.apply(html.unescape)
```

/usr/local/lib/python3.6/dist-packages/pandas/core/generic.py:5303: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
self[name] = value
```

In [ ]:

```

tokenizer = XLNetTokenizer.from_pretrained('xlnet-base-cased')
MAX_SEQUENCE_LENGTH = 512

# this function trims the tokens with length > 512 to match with the bert input.
'''
In the function below, if the input sentence has the number of tokens > 512, the
sentence is trimmed down to 512. To trim the number of tokens, 256 tokens from
the start and 256 tokens from the end are kept and the remaining tokens are dropped.
Ex. suppose an answer has 700 tokens, to trim this down to 512, 256 tokens from the
beginning are taken and 256 tokens from the end are taken and concatenated to make
512 tokens. The remaining [700-(256+256) = 288] tokens that are in the middle of the
answer are dropped. The logic makes sense because in large texts, the beginning part
usually describes what the text is all about and the end part describes the conclusion
of the text. This is also closely related to the target features that we need to predict.
'''
def _trim_input(question_tokens, answer_tokens, max_sequence_length=512, q_max_len=254, a_max_len=254):
    q_len = len(question_tokens)
    a_len = len(answer_tokens)
    if q_len + a_len + 3 > max_sequence_length:
        if a_max_len <= a_len and q_max_len <= q_len:
            q_new_len_head = q_max_len//2
            question_tokens = question_tokens[:q_new_len_head] + question_tokens[-q_new_len_head:]
            a_new_len_head = a_max_len//2
            answer_tokens = answer_tokens[:a_new_len_head] + answer_tokens[-a_new_len_head:]
        elif q_len <= a_len and q_len < q_max_len:
            a_max_len = a_max_len + (q_max_len - q_len - 1)
            a_new_len_head = a_max_len//2
            answer_tokens = answer_tokens[:a_new_len_head] + answer_tokens[-a_new_len_head:]
        elif a_len < q_len:
            q_max_len = q_max_len + (a_max_len - a_len - 1)
            q_new_len_head = q_max_len//2
            question_tokens = question_tokens[:q_new_len_head] + question_tokens[-q_new_len_head:]

    return question_tokens, answer_tokens

```



In [ ]:

```

# function for tokenizing the input data for bert
def _convert_to_transformer_inputs(title, question, answer, tokenizer, question_only=False):
    question = f"{title} [SEP] {question}"
    question_tokens = tokenizer.tokenize(question)
    if question_only:
        answer_tokens = []
    else:
        answer_tokens = tokenizer.tokenize(answer)
    question_tokens, answer_tokens = _trim_input(question_tokens, answer_tokens)
    ids = tokenizer.convert_tokens_to_ids(["[CLS]"] + question_tokens + ["[SEP]"] + answer_tokens + ["[SEP]"])
    padded_ids = ids + [tokenizer.pad_token_id] * (MAX_SEQUENCE_LENGTH - len(ids))
    token_type_ids = [0] * (1 + len(question_tokens) + 1) + [1] * (len(answer_tokens) + 1) + [0] * (MAX_SEQUENCE_LENGTH - len(ids))
    attention_mask = [1] * len(ids) + [0] * (MAX_SEQUENCE_LENGTH - len(ids))
    return padded_ids, token_type_ids, attention_mask

```

In [ ]:

```

# function for creating the input_ids, masks and segments for the bert input
def compute_input_arrays(df, question_only=False):
    input_ids, input_token_type_ids, input_attention_masks = [], [], []
    for title, body, answer in zip(df["question_title"].values, df["question_body"].values, df["answer"].values):
        ids, type_ids, mask = _convert_to_transformer_inputs(title, body, answer, tokenizer, question_only=question_only)
        input_ids.append(ids)
        input_token_type_ids.append(type_ids)
        input_attention_masks.append(mask)
    return (
        np.asarray(input_ids, dtype=np.int32),
        np.asarray(input_attention_masks, dtype=np.int32),
        np.asarray(input_token_type_ids, dtype=np.int32)
    )

def compute_output_arrays(df):
    return np.asarray(df[output_categories])

```

In [ ]:

```
# Creating the model
K.clear_session()
max_seq_length = 512

input_tokens = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="input_tokens")
input_mask = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="input_mask")
# input_segment = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="input_segment")

#bert layer
xlnet_config = XLNetConfig.from_pretrained('xlnet-base-cased', output_hidden_states=True)
xlnet_model = TFXLNetModel.from_pretrained('xlnet-base-cased', config=xlnet_config)

sequence_output, hidden_states = xlnet_model([input_tokens, input_mask])

# Last 4 hidden layers of bert
h12 = tf.reshape(hidden_states[-1][:,0], (-1,1,768))
h11 = tf.reshape(hidden_states[-2][:,0], (-1,1,768))
h10 = tf.reshape(hidden_states[-3][:,0], (-1,1,768))
h09 = tf.reshape(hidden_states[-4][:,0], (-1,1,768))
concat_hidden = tf.keras.layers.Concatenate(axis=2)([h12, h11, h10, h09])

x = GlobalAveragePooling1D()(concat_hidden)

x = Dropout(0.2)(x)

output = Dense(30, activation='sigmoid')(x)

model_qa = Model(inputs=[input_tokens, input_mask], outputs=output)
```

In [ ]:

```
model_qa.summary()
```

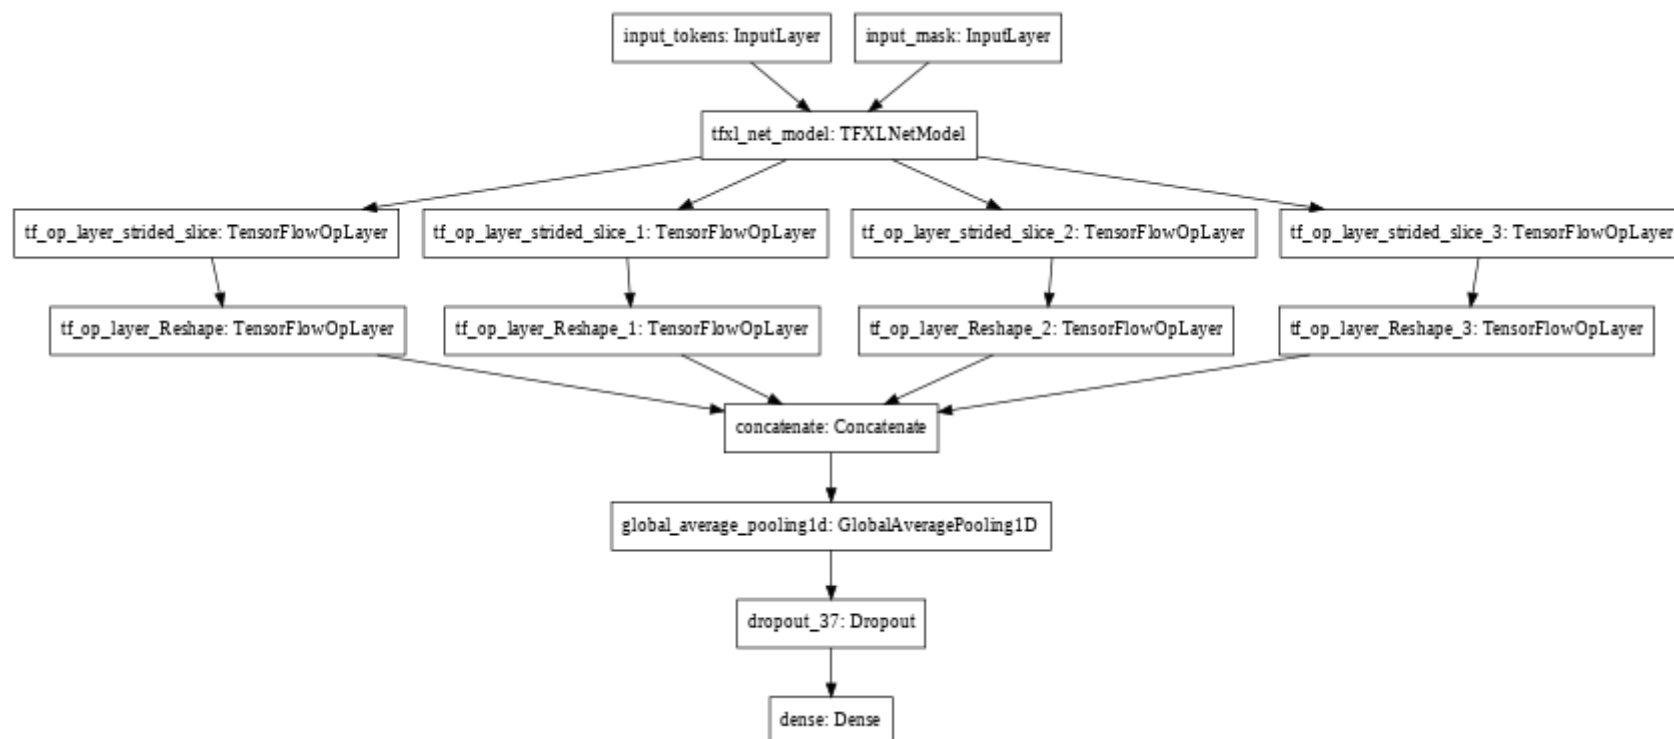
Model: "model"

| Layer (type)                    | Output Shape         | Param #   | Connected to   |
|---------------------------------|----------------------|-----------|--|
| input_tokens (InputLayer)       | [(None, 512)]        | 0         |  |
| input_mask (InputLayer)         | [(None, 512)]        | 0         |  |
| tfxl_net_model (TFXLNetModel)   | ((None, 512, 768), ( | 116718336 | input_tokens[0][0]<br>input_mask[0][0]   |
| tf_op_layer_strided_slice (Tens | [(None, 768)]        | 0         | tfxl_net_model[0][13]  |
| tf_op_layer_strided_slice_1 (Te | [(None, 768)]        | 0         | tfxl_net_model[0][12]  |
| tf_op_layer_strided_slice_2 (Te | [(None, 768)]        | 0         | tfxl_net_model[0][11]  |
| tf_op_layer_strided_slice_3 (Te | [(None, 768)]        | 0         | tfxl_net_model[0][10]  |
| tf_op_layer_Reshape (TensorFlow | [(None, 1, 768)]     | 0         | tf_op_layer_strided_slice[0][0]  |
| tf_op_layer_Reshape_1 (TensorFl | [(None, 1, 768)]     | 0         | tf_op_layer_strided_slice_1[0][0]  |
| tf_op_layer_Reshape_2 (TensorFl | [(None, 1, 768)]     | 0         | tf_op_layer_strided_slice_2[0][0]  |
| tf_op_layer_Reshape_3 (TensorFl | [(None, 1, 768)]     | 0         | tf_op_layer_strided_slice_3[0][0]  |
| concatenate (Concatenate)       | (None, 1, 3072)      | 0         | tf_op_layer_Reshape[0][0]<br>tf_op_layer_Reshape_1[0][0]<br>tf_op_layer_Reshape_2[0][0]<br>tf_op_layer_Reshape_3[0][0] |
| global_average_pooling1d (Globa | (None, 3072)         | 0         | concatenate[0][0]  |
| dropout_37 (Dropout)            | (None, 3072)         | 0         | global_average_pooling1d[0][0]   |
| dense (Dense)                   | (None, 30)           | 92190     | dropout_37[0][0]   |
| Total params: 116,810,526       |                      |           |  |
| Trainable params: 116,810,526   |                      |           |  |
| Non-trainable params: 0         |                      |           |  |

In [ ]:

```
tf.keras.utils.plot_model(
    model_qa, to_file='model.png',
    show_shapes=False,
    show_layer_names=True,
    rankdir='TB',
    expand_nested=False, dpi=48
)
```

Out[ ]:



In [ ]:

```
# test data
tokens, masks, segments = compute_input_arrays(X_test)
test_data = {'input_tokens': tokens,
             'input_mask': masks}
```

In [ ]:

```
# Train data
tokens, masks, segments = compute_input_arrays(X)
def generate_data(tr, cv):
    train_data = {'input_tokens': tokens[tr],
                  'input_mask': masks[tr]}

    cv_data = {'input_tokens': tokens[cv],
               'input_mask': masks[cv]}

    return train_data, cv_data, y.values[tr], y.values[cv]
```

In [ ]:

```
# Function to calculate the Spearman's rank correlation coefficient 'rhos' of actual and predicted data.
from scipy.stats import spearmanr
def compute_spearmanr_ignore_nan(trues, preds):
    rhos = []
    for tcol, pcol in zip(np.transpose(trues), np.transpose(preds)):
        rhos.append(spearmanr(tcol, pcol).correlation)
    return np.nanmean(rhos)
```

In [ ]:

```
# Making the 'rhos' metric to tensorflow graph compatible.
def rhos(y, y_pred):
    return tf.py_function(compute_spearmanr_ignore_nan, (y, y_pred), tf.double)
metrics = [rhos]
```

In [20]:

```
from sklearn.model_selection import KFold
# Compiling and training the model
optimizer = tf.keras.optimizers.Adam(learning_rate=0.00002)
model_qa.compile(loss='binary_crossentropy', optimizer=optimizer, metrics=metrics)
kf = KFold(n_splits=5, random_state=42)
for tr, cv in kf.split(np.arange(train.shape[0])):
    tr_data, cv_data, y_tr, y_cv = generate_data(tr, cv)
    model_qa.fit(tr_data, y_tr, epochs=1, batch_size=4, validation_data=(cv_data, y_cv))
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:296: FutureWarning: Setting a random_state has no effect since shuffle is False. This will raise an error in 0.24. You should leave random_state to its default (None), or set shuffle=True.
```

```
FutureWarning
```



```
WARNING:tensorflow:Gradients do not exist for variables ['tfxl_net_model/transformer/mask_emb:0', 'tfxl_net_model/transformer/layer_.0/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.0/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.1/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.1/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.2/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.2/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.3/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.3/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.4/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.4/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.5/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.5/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.6/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.6/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.7/rel_attn/r_s_bias:0', 'tfxl net model/transformer/layer . 7/rel attn/seg embed:0', 'tfxl net model/transformer/laye
```

```
r_.8/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.8/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.9/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.9/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.10/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.10/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.11/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.11/rel_attn/seg_embed:0'] when minimizing the loss.
```

```
WARNING:tensorflow:Gradients do not exist for variables ['tfxl_net_model/transformer/mask_emb:0', 'tfxl_net_model/transformer/layer_.0/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.0/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.1/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.1/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.2/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.2/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.3/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.3/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.4/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.4/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.5/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.5/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.6/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.6/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.7/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.7/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.8/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.8/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.9/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.9/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.10/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.10/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.11/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.11/rel_attn/seg_embed:0'] when minimizing the loss.
```

```
/usr/local/lib/python3.6/dist-packages/numpy/lib/function_base.py:2534: RuntimeWarning: invalid value encountered in true_divide
```

```
c /= stddev[:, None]
```

```
/usr/local/lib/python3.6/dist-packages/numpy/lib/function_base.py:2535: RuntimeWarning: invalid value encountered in true_divide
```

```
c /= stddev[None, :]
```

```
/usr/local/lib/python3.6/dist-packages/scipy/stats/_distn_infrastructure.py:903: RuntimeWarning: invalid value encountered in greater
```

```
return (a < x) & (x < b)
```

```
/usr/local/lib/python3.6/dist-packages/scipy/stats/_distn_infrastructure.py:903: RuntimeWarning: invalid value encountered in less
```

```
return (a < x) & (x < b)
```

```
/usr/local/lib/python3.6/dist-packages/scipy/stats/_distn_infrastructure.py:1912: RuntimeWarning: invalid value encountered in less_equal
```

```
cond2 = cond0 & (x <= _a)
```

```
1216/1216 [=====] - 2497s 2s/step - loss: 0.4874 - rhos: 0.1715 - val_loss:
0.3885 - val_rhos: 0.3570
1216/1216 [=====] - 2498s 2s/step - loss: 0.4249 - rhos: 0.2746 - val_loss:
0.3726 - val_rhos: 0.4055
1216/1216 [=====] - 2497s 2s/step - loss: 0.4050 - rhos: 0.3183 - val_loss:
0.3647 - val_rhos: 0.4501
1216/1216 [=====] - 2498s 2s/step - loss: 0.3899 - rhos: 0.3510 - val_loss:
0.3483 - val_rhos: 0.4768
WARNING:tensorflow:Gradients do not exist for variables ['tfxl_net_model/transformer/mask_emb:0', 't
fxl_net_model/transformer/layer_.0/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.0/rel_
attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.1/rel_attn/r_s_bias:0', 'tfxl_net_model/trans
former/layer_.1/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.2/rel_attn/r_s_bias:0',
'tfxl_net_model/transformer/layer_.2/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.3/r
el_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.3/rel_attn/seg_embed:0', 'tfxl_net_model/tr
ansformer/layer_.4/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.4/rel_attn/seg_embed:
0', 'tfxl_net_model/transformer/layer_.5/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.
5/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.6/rel_attn/r_s_bias:0', 'tfxl_net_mode
l/transformer/layer_.6/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.7/rel_attn/r_s_bi
as:0', 'tfxl_net_model/transformer/layer_.7/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/laye
r_.8/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.8/rel_attn/seg_embed:0', 'tfxl_net_m
odel/transformer/layer_.9/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.9/rel_attn/seg_
embed:0', 'tfxl_net_model/transformer/layer_.10/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/l
ayer_.10/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.11/rel_attn/r_s_bias:0', 'tfxl_
net_model/transformer/layer_.11/rel_attn/seg_embed:0'] when minimizing the loss.
WARNING:tensorflow:Gradients do not exist for variables ['tfxl_net_model/transformer/mask_emb:0', 't
fxl_net_model/transformer/layer_.0/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.0/rel_
attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.1/rel_attn/r_s_bias:0', 'tfxl_net_model/trans
former/layer_.1/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.2/rel_attn/r_s_bias:0',
'tfxl_net_model/transformer/layer_.2/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.3/r
el_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.3/rel_attn/seg_embed:0', 'tfxl_net_model/tr
ansformer/layer_.4/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.4/rel_attn/seg_embed:
0', 'tfxl_net_model/transformer/layer_.5/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.
5/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.6/rel_attn/r_s_bias:0', 'tfxl_net_mode
l/transformer/layer_.6/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.7/rel_attn/r_s_bi
as:0', 'tfxl_net_model/transformer/layer_.7/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/laye
r_.8/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.8/rel_attn/seg_embed:0', 'tfxl_net_m
odel/transformer/layer_.9/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/layer_.9/rel_attn/seg_
embed:0', 'tfxl_net_model/transformer/layer_.10/rel_attn/r_s_bias:0', 'tfxl_net_model/transformer/l
ayer_.10/rel_attn/seg_embed:0', 'tfxl_net_model/transformer/layer_.11/rel_attn/r_s_bias:0', 'tfxl_
net_model/transformer/layer_.11/rel_attn/seg_embed:0'] when minimizing the loss.
1216/1216 [=====] - 2506s 2s/step - loss: 0.3721 - rhos: 0.3942 - val_loss:
0.3452 - val_rhos: 0.5013
```

In [27]:

```
model_qa.save_weights("drive/My Drive/xlnet_model_qa.h5")
```

In [23]:

```
# train data
tokens, masks, segments = compute_input_arrays(X)
train_data = {'input_tokens': tokens,
              'input_mask': masks}
```

In [25]:

```
# Predicting the train and test data labels
pred_a_test = model_qa.predict(test_data)
pred_a_train = model_qa.predict(train_data)

# saving the predicted labels as dataframes
df = pd.DataFrame(pred_a_train, columns=y.columns)
df.to_csv('xlnet_pred_qa_train.csv', index=False)

df = pd.DataFrame(pred_a_test, columns=y.columns)
df.to_csv('xlnet_pred_qa_test.csv', index=False)
```

In [ ]: